

Mark Do

+1 647-766-6794 | mdo@uwaterloo.com | github.com/hepromark | linkedin.com/in/markdouw/

Education

University of Waterloo - Candidate for *Bachelor of Mechatronics Engineering*

2022 - 2027

- 4.0 Cumulative GPA | 94% term average | 3 x Dean's Honours List

Work Experience

Algorithms Software Developer | *Python, SimPy*

Sep 2024 – Dec 2024

Rapyuta Robotics

Tokyo, Japan

- Created an end-to-end pipeline that generates optimized Robotic Storage System layouts based on customer requirements.
- Developed a discrete-event simulator using spatial-temporal grids to estimate throughput for combinations of different item distributions, path-planning constraints & scheduling algorithms for feasibility testing of 5+ new system improvements.
- Utilized Genetic Algorithms to do non-convex optimization, resulting in 10% higher storage density without speed penalties.

Robotics Software Developer | *Python, Django, SQL, ROS*

May 2024 – Sep 2024

Rapyuta Robotics

Tokyo, Japan

- Developed a Task Planner component to reactively schedule tasks for robots in a Multi-Agent Robotic Storage System, which translates customer warehouse operations into task trees appropriate for parallelized robot execution.
- Scaled Task Planner & world model database to work with 10 times more robots (6 → 60) by using caches & lighter queries.
- Raised reliability significantly by creating 5+ error recovery behaviours, resulting in system MTBF gains from hours to days.

Test Automation Developer | *Java, Docker, TestNG*

Sep 2023 – Dec 2023

Definity Financial

Waterloo, Canada

- Overhauled Java test suite with > 120 test cases for company's mobile app, reducing false positive bug detections by > 95%.
- Developed a new automated module to dynamically create test batches & execute them in parallel to reduce testing time.
- Created a new automated CI pipeline for regression tests, executing & reporting results daily to save >12 dev hrs./week.

Related Experience

Autonomous Racing Team Lead | *C++, Docker, ROS2*

Jan 2025 - Present

F1Tenth Autonomous Racing Team

- Leading a team of 15+ undergraduates to develop a self-driving software stack for F1Tenth Autonomous Racing Competition.
- Led the design & implementation of pure pursuit control, RRT motion planning, offline & online SLAM algorithms.
- Built a dockerized dev environment integrating ROS2, Foxglove and AutoDRIVE physics simulator for testing algorithms.

Autonomous Vehicles Developer | *C++, Casadi, Kubernetes, Terraform, PyTree*

Nov 2023 – Jul 2024

WATONOMOUS (Waterloo Autonomous Vehicles Design Team)

- Developed a dual-model MPC system, joining a Multi-Layer Perceptron with a traditional kinematics model for car navigation.
- Designed a high-level car navigation controller that uses behaviour trees to take driving actions based on GPS data.
- Maintained SLURM & SSH server infrastructure for over 400 developers across 3 universities for ML research & development.
- Created an Asset Manager that manages website assets using self-hosted Ceph S3 Buckets & deployed on a Kubernetes Cluster.

Autonomy Developer | *Unsupervised Learning, Scikit-learn, Linux, Clustering algorithms*

Apr 2023 – Sep 2023

Waterloo Aerial Robotics Group

- Developed a [clustering algorithm](#) for an autonomous drone's onboard perception system to postprocess landing pad CV model.
- Implemented a VGMM as the clustering algorithm to differentiate real landing locations from false positives.

Technical Skills

Languages: Python, C, C++, Java, JavaScript, Bash, Arduino, MATLAB, Assembly


Libraries & Frameworks: ROS, Gazebo, Foxglove, RVIZ, PyTorch, OpenCV, Scikit-learn, PyTree

DevOps: Docker, Kubernetes, Terraform, Ansible, Git, GitHub, BitBucket, AWS

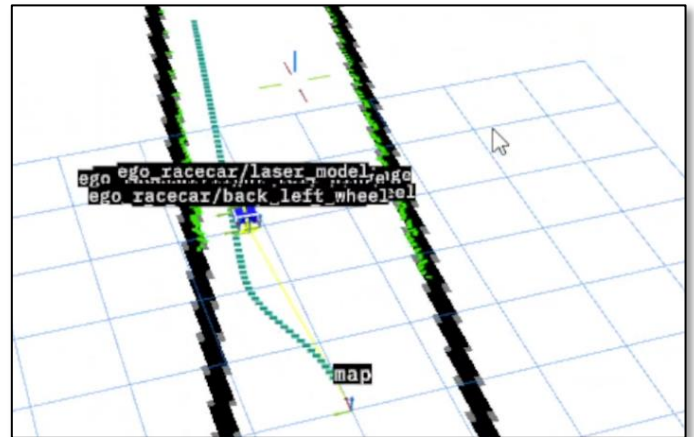
Mark Do - Portfolio

647-394-6627 | mdu@uwaterloo.com | github.com/hepromark | linkedin.com/in/markdouw/

F1Tenth Autonomous Racing Team Lead



F1Tenth Vehicle Simulator



Pure-pursuit control applied to a spline generated by RRT

The F1Tenth competition is a 1/10th scale autonomous racing cup where global teams compete for the fastest autonomy algorithms.


- Leading a team of 15+ undergraduates to develop a self-driving software stack for F1Tenth Autonomous Racing Competition.
- Modelled the onboard software architecture after the See-Process-Act paradigm, prioritizing real-time control & safety.
- Led the design & implementation of pure pursuit control, RRT motion planning, offline & online SLAM algorithms.
- Built a dockerized dev environment integrating ROS2, Foxglove and AutoDRIVE physics simulator.
- Project Repo can be found [here](#).



Landing Pad Clustering Algorithm: In-flight unsupervised learning model



WARG's 2023 Competition Drone: ICARUS


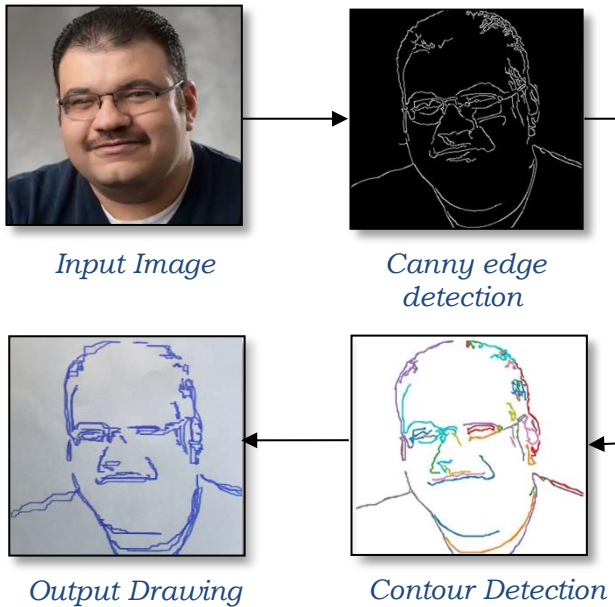


Clustering Algorithm Output: False positives are discarded from potential landing pads.

I worked on WARG's in-flight unsupervised learning model for the 2023 AEAC Competition

- Developed a clustering module for an autonomous drone to cluster landing pad detections & predict actual pad locations.
- Implemented a Variational Gaussian mixture model to consume pad locations over time and continuously output estimates.
- Created filters with dynamic thresholding to remove outlier pad locations, enabling drone to operate in poor visibility weather.
- Designed an intelligent memory system to bias towards new observations while retaining knowledge of pads only seen once.
- Code can be found [here](#)

Articus Maximus: Sketching Robot



Sketching Robot after completing a drawing

- Designed & manufactured a 2-axis gantry sketching robot which controls a pen to draw images on paper from digital file input.
- Developed an image pipeline that extracts drawable contours from digital images with edge detection & contour detection.
- Optimized robot drawing speed by using the Douglas Peucker algorithm and Hu Moments to reduce total contour count.
- Programmed and tuned a PID controller in C with anti-windup & a 1D motion profile to draw lines accurate within 2 degrees.
- Code can be found [here](#).

BlindWatchers: Assistive Vision Headwear



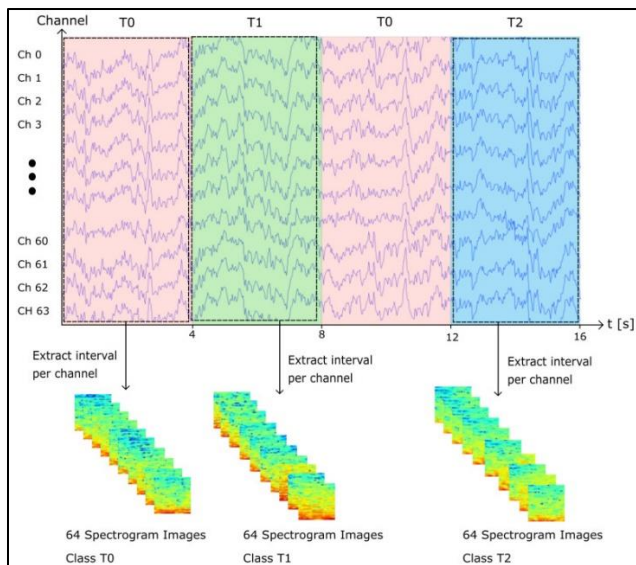
Headwear has 2 cameras, headphones for audio output, with an NVIDIA Jetson attached behind for compute.




Sample YoloV8 detection bounding boxes, before conversion to spatial audio.

- Designed and developed a smart helmet for the visually impaired with an NVIDIA Jetson, integrating YOLOv8 object detection, directional audio, & Google Cloud Speech-to-Text to provide real-time auditory descriptions of nearby objects.
- Estimated object real-world 3D poses from 2D bounding boxes to provide matching directional audio for more intuitive UX.
- Implemented an asynchronous architecture with a fast-inferencing CV model to decrease latency between detection & audio.
- Code can be found [here](#)

EEG Motor Imagery Classifier




CNN Architecture takes 3D input volumes, created from stacking 64 electrodes' spectrogram images



Spectrogram image produced from a single electrode's 2-minute EEG segment

- Developing deep learning model on EEG brain data for motor imagery classification task on a 64-channel Motor Imagery dataset.
- Converted 64-channel time series data into volumes in the time-frequency domain using FFTs for use with convolutional kernels.
- Wrote custom data samplers in PyTorch with weak shuffling to optimize data read speeds for training batches from .h5 files.
- Project Repo can be found [here](#)

Soil Humidity IoT: Cloud-based agriculture monitoring.



Examples of soil-humidity detector ESP8266 setup




Diagram of the agriculture monitoring system.

- Designed and created a monitoring system to track houseplant soil humidity, alerting users via SMS when below set threshold
- Programmed ESP8266s to publish humidity levels via Wi-Fi to an AWS IoT MQTT topic, storing the data in Dynamo DB.
- Created an AWS Lambda function listening to database writes, to alert users using AWS SNS if readings below threshold.
- Code can be found [here](#)